

Query-aware pruning: answer accuracy vs. budget

